

The role of Artificial Intelligence in future technology

Andreas Giannoutsos

May 16, 2022

Abstract

Artificial intelligence (AI) has evolved significantly in recent years. The goal of this study is to understand current achievements and estimate their influence on future technologies. In this investigation, we begin by looking back at the historical evolution of artificial intelligence. Then, after completing an examination of its present accomplishments in numerous domains, we will clarify the impact it will have on future technologies. Finally, we address its limitations, the challenges it faces, and the concerns it raises.

1 Introduction

Artificial intelligence (AI) is intelligence exhibited by computational machines. Artificial intelligence research is the study of intelligent behavior and the way we seek to replicate and develop it by using computer equipment. It is also used to describe any computer that has human-like traits such as learning and problem-solving.

It is a broad term that refers to a variety of methodologies and algorithms that aim to create systems that simulate ways of thinking. It has been employed in many fields of computer science but also in other industries and businesses such as the automotive, aerospace, medical, and robotics industries.

Artificial intelligence has made significant advances in recent years, notably in the machine learning sub-field known as deep learning. AI is already supporting people in their daily lives and has the potential to assist them even further by contributing to the creation of new technological tools. However, despite of this rapid progress, AI development has already revealed limitations. Concerns have also been raised about the possible uses of these advancements, as well as their impact on future technologies.

2 The Origins of AI

For a long time, people have imagined objects coming to life as sentient entities. Many ancient civilizations believed robots were mythical. The origins of current artificial intelligence may be traced back to classical philosophers' attempts to classify human thought as a symbolic system.

In our modern times, science fiction introduced the world to the notion of artificially intelligent robots in the first part of the twentieth century. By the 1950s, a generation of scientists, mathematicians, and philosophers were culturally assimilated to the notion of artificial intelligence. Alan Turing, investigated the math-

ematical possibilities of artificial intelligence.

Turing proposed that humans use available information as well as reason to solve problems and make decisions, so why can't machines? This was the conceptual basis for his 1950 study [63], in which he addressed how to develop intelligent machines and measure their intelligence with the famous Turing Test. The Turing Test was created to establish an adequate operational definition of intelligence. A computer passes the test if a human interrogator cannot distinguish whether the written replies are human made or from a computer. Natural language processing, knowledge representation, automated reasoning, machine learning, computer vision, and robotics are just a few of the tests that the AI system must satisfy.

However, the area of artificial intelligence was not formally established until 1956, when the phrase "artificial intelligence" was invented [40].

AI thrived from the 1950s until the mid '70s. Machine learning algorithms improved as well, and individuals became more adept at determining which method to apply to a given situation. Early demonstrations, like ELIZA, showed promise in terms of problem solving goals [67]. However, issues emerged since most early programs understood little about their subject area. The interest in AI was sparked again in the 1980s by the growth of algorithmic tools and computer processing capabilities. Researchers reinvented the back-propagation algorithm. Geoffrey Hinton, John Hopfield, and David Rumelhart pioneered "deep learning" techniques, which allowed computers to learn based on their own experiences [27] [51].

Deep learning is a technique that enables computers to learn abstract data representations through many layers of mathematical operations and parameters [34]. These methods have dramatically improved the state-of-the-art in speech recognition, visual object recognition, object detection, and many other domains such as drug discovery and genomics in today's "big data"

age.

3 Applications of AI and their and Future Impact

Despite its basis in computer science, AI has become a critical component in a variety of industries. Initially, AI is used in the medicine and bio-informatics industries. The detection and diagnosis [1] in medical images using deep learning approaches in 2d [12] [49] and even 3d [14] can be made with accuracy comparable to that of a doctor. AI will not replace doctors, but it will help them to improve their everyday medical processes and hence influence positively their work. Specialities which have substantial visual elements, will be significantly shaped [20]. Furthermore, important advances in the science of protein synthesis have been made, as a result of the AlphaFold model [55] [30]. These have resulted in enormous discoveries in structural and molecular biology, which can aid in the synthesis of proteins for a number of treatments [50]. These discoveries bring us one step closer to simulating the operation of the human body, which will allow scientists to develop a framework for combating even the most challenging diseases.

AI has transformed both the fields of art and music with generative models that can create new unique pieces of art [18] [29] [31]. Research in multi-modal models has resulted in models that may produce visuals from text, offering designers new options the combination of unrelated contexts [47] [46].

In addition, AI has made significant advances in the field of Natural Language Processing. From the period of LSTMs [26] to the era of Transformers [64], language models have progressed at an unparalleled rate, reaching human-level performance in many circumstances. The development of massive language models [17] [37] [45] [9] has revolutionized numerous domains, including machine translation [69], zero-shot learning [62], question and answering [70], text classifications [61] and summarization [2]. These language models are also trained to generate semantic data, they can write code that outperforms coding competitions [13], and solve mathematical problems that outperform mathematical competitions [15] too. These models have the potential to greatly improve the professions of developers and scientists in the future and enhance their productivity to greater levels.

In the realm of Reinforcement Learning, AI has taken a major step. This type of AI provides the ability to teach robotic agents to perform new complicated processes without the need for pre-existing data [4] [22]. This will result in the development of robots that will learn from their environment and assist people in many aspects of their lives. With the introduction of new innovative algorithms like [43] and [54] more and more Reinforcement Learning systems manage to outperform humans in incredibly complex operations.

From the AlphaZero system [56] that beat humans in the already complex game of chess, to other AI systems that beat human masters in the games of Go [57], Dota [6] and Starcraft [65] which are more complicated and demand greater intelligence and long term strategic skills, machines outperform humans in tasks that they never imagined.

Moreover, AI has benefited the field of computer vision. Due to the complexity of the optical process, the abundance of data and the recent deep learning approaches have facilitated major strides in computer vision challenges. In the area of image classification, machines have reached the level of humans [52], while significant progress has been made in object detection [21], image segmentation [38], and object tracking [3]. Following that work, the Dynamic Vision and Learning Group laboratory at the Technical University of Munich is aiming to improve the evaluation metrics of object tracking [36] [66] while at the same time trying to increase them. The TrackFormer model [42] which predicts multiple objects in space and time in an end-to-end manner, outperforms more complicated models that rely on motion and graph models. These advancements, will accelerate the productivity and the safety of workers, and in turn, they have the potential to be applied to a wide range of fields, like in transportation with self-driving cars [11] and pedestrian tracking [25], in manufacturing with quality management [68], in agriculture with yield assessment [32] and in many others sectors.

Aside from computer vision, AI has been used in physics to accelerate the modeling of physical phenomena such as fluid and physics simulations [53] [33] [39], as well as weather forecasting [48]. It has also been used to answer fusion reaction mysteries [16] [19], to analyze particle accelerator discoveries [23] and to enhance the entanglement capacity of quantum systems [35]. The value that AI will provide through facilitating scientific study in the field of physics is immeasurable, as most of today's technology feats stem on scientific discoveries made in the past century.

Finally, the combination of these new ideas, methods and people together will contribute to various prospective AI applications that will have a significant influence on future technology [24].

4 Limitations and Concerns

While breakthroughs in artificial intelligence research are projected to bring significant and innovative applications in the future, the technology may be less revolutionary in the long run than many anticipate.

Although artificial intelligence may outperform humans in complicated activities such as video games and chess, it cannot replicate human behaviour. Because of this contradiction, researchers have classified AI into two types: artificial general intelligence (AGI)

or strong AI [44] and weak or narrow AI. Strong AI is able to learn any task that humans can execute. Weak AI, on the other hand, is not supposed to have cognitive capabilities; it is a system designed to tackle a specific issue.

The evolution of strong AI is considered particularly difficult as many factors that contribute to biological intelligence and their incorporation into the artificially intelligent one are not yet understood. Intuitive physics is the ability of humans to grasp how things interact with one another. Humans have an intuitive understanding of physics that helps them to traverse the environment and make predictions, even in completely new situations. AI, on the other hand, is currently incapable of replicating this behavior. Furthermore, modern AI systems lack intuitive psychology, which is the ability to obtain insight into the motivations of sentient creatures and make predictions based on them. This intuition enables individuals to foresee and prepare their reactions to the behaviors of others. AI struggles in this area too. These particular structures in the human brain allow this subconscious decision-making. Researchers are struggling to replicate the adaptive behavior process in AI systems since the determination of terminal loss or reward function for such dynamic processes is a very hard task [58]. Therefore, in the future, the development of the Singularity [10] is quite unlikely.

On the other hand, the development of narrow AI, guided by recent innovations, discoveries and progress in recent times, seems quite promising. Advanced programming, mathematical and data-engineering techniques will result in more sophisticated single-task accomplishments. This type of AI has applications in many industries, but it has limitations since it cannot perform another function after it has been devoted to one. Nonetheless, weak AI could free up human labor capacities and allow for more accurate work. It may also help to reduce expenses and improve task efficiencies.

While artificial general intelligence is highly unlikely to be developed in the near future, it also remains a major societal worry. AGI might supplement human intelligence by providing a greater capacity to address the world's issues. Alternatively, whoever owns it might use it maliciously.

Nonetheless, even today, while being far from AGI, phenomena like the Cambridge Analytica scandal [28] [7] have been witnessed in which AI was utilized to drastically impact public opinion. Many research have also found that these AI models discriminate between social groups of humans [41] [5] [8]. Furthermore, the present large deep learning models, as well as the increased computer power required for their training, have a high carbon footprint and negatively contribute to climate change [59] [60]. As a result, various improvements and upgrades to the models' performance and energy requirements are necessary. It is also vi-

tal to create better methods to avoid bias rather than simply filtering the discriminatory data generated by people.

In such instances, the outcome might be determined by the AGI's goals: whether it seeks something beneficial for humans, the planet, itself, or something altogether else.

5 Conclusion

The origins and evolution of AI can teach us a lot about future advancement. The progression of AI nowadays is triggered by the increase in computational capacity as well as the abundance of data. However, there are limitations to this remarkable breakthrough, as the nature of deep learning mathematical models makes AGI development challenging. Simultaneously, there are concerns regarding the misuse of AI. As a result, appropriate financing and a broader strategy focusing on AI research and development based on ethics and the enhancement of people's lives are essential. As with any technical invention, it is up to people to decide if it will be utilized for positive or negative purposes.

References

- [1] Ravi Aggarwal et al. "Diagnostic accuracy of deep learning in medical imaging: a systematic review and meta-analysis". In: *npj Digital Medicine* 4 (Dec. 2021). DOI: 10.1038/s41746-021-00438-z.
- [2] Armen Aghajanyan et al. "Better Fine-Tuning by Reducing Representational Collapse". In: *International Conference on Learning Representations*. 2021. URL: <https://openreview.net/forum?id=0Q08SN70M1V>.
- [3] Alexandre Alahi et al. "Social LSTM: Human Trajectory Prediction in Crowded Spaces". In: *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2016, pp. 961–971. DOI: 10.1109/CVPR.2016.110.
- [4] OpenAI: Marcin Andrychowicz et al. "Learning dexterous in-hand manipulation". In: *The International Journal of Robotics Research* 39.1 (2020), pp. 3–20. DOI: 10.1177/0278364919887447. eprint: <https://doi.org/10.1177/0278364919887447>. URL: <https://doi.org/10.1177/0278364919887447>.
- [5] Emily M. Bender et al. "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?" In: *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*. FAccT '21. Virtual Event, Canada: Association for Computing Machinery, 2021, pp. 610–623. ISBN: 9781450383097. DOI: 10.1145/3442188.3445922. URL: <https://doi.org/10.1145/3442188.3445922>.
- [6] Christopher Berner et al. "Dota 2 with Large Scale Deep Reinforcement Learning". In: *ArXiv abs/1912.06680* (2019).
- [7] Elena Boldyreva. "Cambridge Analytica: Ethics And Online Manipulation With Decision-Making Process". In: Dec. 2018, pp. 91–102. DOI: 10.15405/epsbs.2018.12.02.10.

- [8] Tolga Bolukbasi et al. “Man is to Computer Programmer as Woman is to Homemaker? Debiasing Word Embeddings”. In: *Proceedings of the 30th International Conference on Neural Information Processing Systems. NIPS’16*. Barcelona, Spain: Curran Associates Inc., 2016, pp. 4356–4364. ISBN: 9781510838819.
- [9] Tom Brown et al. “Language Models are Few-Shot Learners”. In: *Advances in Neural Information Processing Systems*. Ed. by H. Larochelle et al. Vol. 33. Curran Associates, Inc., 2020, pp. 1877–1901. URL: <https://proceedings.neurips.cc/paper/2020/file/1457c0d6bfc4967418bfb8ac142f64a-Paper.pdf>.
- [10] David J. Chalmers. “The Singularity: A Philosophical Analysis”. In: *Journal of Consciousness Studies* 17.9-10 (2010), pp. 9–10.
- [11] Dian Chen and Philipp Krähenbühl. “Learning from all vehicles”. In: *CVPR*. 2022.
- [12] Liang-Chieh Chen et al. “DeepLab: Semantic Image Segmentation with Deep Convolutional Nets, Atrous Convolution, and Fully Connected CRFs”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 40.4 (2018), pp. 834–848. DOI: 10.1109/TPAMI.2017.2699184.
- [13] Mark Chen et al. *Evaluating Large Language Models Trained on Code*. 2021. DOI: 10.48550/ARXIV.2107.03374. URL: <https://arxiv.org/abs/2107.03374>.
- [14] Özgün Çiçek et al. “3D U-Net: Learning Dense Volumetric Segmentation from Sparse Annotation”. In: *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2016*. Ed. by Sebastien Ourselin et al. Cham: Springer International Publishing, 2016, pp. 424–432. ISBN: 978-3-319-46723-8.
- [15] Karl Cobbe et al. *Training Verifiers to Solve Math Word Problems*. 2021. DOI: 10.48550/ARXIV.2110.14168. URL: <https://arxiv.org/abs/2110.14168>.
- [16] Jonas Degraeve et al. “Magnetic control of tokamak plasmas through deep reinforcement learning”. In: *Nature* 602.7897 (Feb. 2022), pp. 414–419. ISSN: 1476-4687. DOI: 10.1038/s41586-021-04301-9. URL: <https://doi.org/10.1038/s41586-021-04301-9>.
- [17] Jacob Devlin et al. “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding”. In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Minneapolis, Minnesota: Association for Computational Linguistics, June 2019, pp. 4171–4186. DOI: 10.18653/v1/N19-1423. URL: <https://aclanthology.org/N19-1423>.
- [18] Prafulla Dhariwal et al. “Jukebox: A Generative Model for Music”. In: *arXiv preprint arXiv:2005.00341* (2020).
- [19] Diogo R. Ferreira. *Applications of Deep Learning to Nuclear Fusion Research*. 2018. DOI: 10.48550/ARXIV.1811.00333. URL: <https://arxiv.org/abs/1811.00333>.
- [20] A. Fourcade and R.H. Khonsari. “Deep learning in medical image analysis: A third eye for doctors”. In: *Journal of Stomatology, Oral and Maxillofacial Surgery* 120.4 (2019). 55th SFSCMFCO Congress, pp. 279–288. ISSN: 2468-7855. DOI: <https://doi.org/10.1016/j.jormas.2019.06.002>. URL: <https://www.sciencedirect.com/science/article/pii/S2468785519301582>.
- [21] Ross Girshick. “Fast R-CNN”. In: *2015 IEEE International Conference on Computer Vision (ICCV)*. 2015, pp. 1440–1448. DOI: 10.1109/ICCV.2015.169.
- [22] Shixiang Gu et al. “Deep reinforcement learning for robotic manipulation with asynchronous off-policy updates”. In: *2017 IEEE International Conference on Robotics and Automation (ICRA)*. 2017, pp. 3389–3396. DOI: 10.1109/ICRA.2017.7989385.
- [23] Dan Guest, Kyle Cranmer, and Daniel Whiteson. “Deep Learning and Its Application to LHC Physics”. In: *Annual Review of Nuclear and Particle Science* 68.1 (2018), pp. 161–181. DOI: 10.1146/annurev-nucl-101917-021019. eprint: <https://doi.org/10.1146/annurev-nucl-101917-021019>. URL: <https://doi.org/10.1146/annurev-nucl-101917-021019>.
- [24] Andrew Hargadon. *How breakthroughs happen: The surprising truth about how companies innovate*. Harvard Business Press, 2003.
- [25] Roberto Henschel et al. “Fusion of Head and Full-Body Detectors for Multi-object Tracking”. In: *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*. 2018, pp. 1509–150909. DOI: 10.1109/CVPRW.2018.00192.
- [26] Sepp Hochreiter and Jürgen Schmidhuber. “Long Short-Term Memory”. In: *Neural Comput.* 9.8 (Nov. 1997), pp. 1735–1780. ISSN: 0899-7667. DOI: 10.1162/neco.1997.9.8.1735. URL: <https://doi.org/10.1162/neco.1997.9.8.1735>.
- [27] J J Hopfield. “Neural networks and physical systems with emergent collective computational abilities.” In: *Proceedings of the National Academy of Sciences* 79.8 (1982), pp. 2554–2558. DOI: 10.1073/pnas.79.8.2554. URL: <https://www.pnas.org/doi/abs/10.1073/pnas.79.8.2554>.
- [28] Margaret Hu. “Cambridge Analytica’s black box”. In: *Big Data & Society* 7.2 (2020), p. 2053951720938091. DOI: 10.1177/2053951720938091. eprint: <https://doi.org/10.1177/2053951720938091>. URL: <https://doi.org/10.1177/2053951720938091>.
- [29] Yongcheng Jing et al. “Neural Style Transfer: A Review”. In: *IEEE Transactions on Visualization and Computer Graphics* 26.11 (2020), pp. 3365–3385. DOI: 10.1109/TVCG.2019.2921336.
- [30] John Jumper et al. “Highly accurate protein structure prediction with AlphaFold”. In: *Nature* 596.7873 (Aug. 2021), pp. 583–589. ISSN: 1476-4687. DOI: 10.1038/s41586-021-03819-2. URL: <https://doi.org/10.1038/s41586-021-03819-2>.
- [31] Tero Karras, Samuli Laine, and Timo Aila. “A Style-Based Generator Architecture for Generative Adversarial Networks”. In: *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2019, pp. 4396–4405. DOI: 10.1109/CVPR.2019.00453.

- [32] Saeed Khaki and Lizhi Wang. “Crop Yield Prediction Using Deep Neural Networks”. In: *Frontiers in Plant Science* 10 (2019). ISSN: 1664-462X. DOI: 10.3389/fpls.2019.00621. URL: <https://www.frontiersin.org/article/10.3389/fpls.2019.00621>.
- [33] Byungsoo Kim et al. “Deep Fluids: A Generative Network for Parameterized Fluid Simulations”. In: *Computer Graphics Forum* 38 (2019).
- [34] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. “Deep learning”. In: *nature* 521.7553 (2015), p. 436.
- [35] Yoav Levine et al. “Quantum Entanglement in Deep Learning Architectures”. In: *Phys. Rev. Lett.* 122 (6 Feb. 2019), p. 065301. DOI: 10.1103/PhysRevLett.122.065301. URL: <https://link.aps.org/doi/10.1103/PhysRevLett.122.065301>.
- [36] Yang Liu et al. “Opening up Open-World Tracking”. In: *CoRR* abs/2104.11221 (2021). arXiv: 2104.11221. URL: <https://arxiv.org/abs/2104.11221>.
- [37] Yinhan Liu et al. *RoBERTa: A Robustly Optimized BERT Pretraining Approach*. 2019. DOI: 10.48550/ARXIV.1907.11692. URL: <https://arxiv.org/abs/1907.11692>.
- [38] J. Long, E. Shelhamer, and T. Darrell. “Fully convolutional networks for semantic segmentation”. In: *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. Los Alamitos, CA, USA: IEEE Computer Society, June 2015, pp. 3431–3440. DOI: 10.1109/CVPR.2015.7298965. URL: <https://doi.ieeecomputersociety.org/10.1109/CVPR.2015.7298965>.
- [39] Michael Lutter, Christian Ritter, and Jan Peters. “Deep Lagrangian Networks: Using Physics as Model Prior for Deep Learning”. In: *International Conference on Learning Representations*. 2019. URL: <https://openreview.net/forum?id=BklHjCqKm>.
- [40] John McCarthy. “Artificial Intelligence, Logic and Formalizing Common Sense”. In: *Philosophical Logic and Artificial Intelligence*. Ed. by Richmond H. Thomason. Dordrecht: Springer Netherlands, 1989, pp. 161–190. ISBN: 978-94-009-2448-2. DOI: 10.1007/978-94-009-2448-2_6. URL: https://doi.org/10.1007/978-94-009-2448-2_6.
- [41] Ninareh Mehrabi et al. “A Survey on Bias and Fairness in Machine Learning”. In: *ACM Comput. Surv.* 54.6 (July 2021). ISSN: 0360-0300. DOI: 10.1145/3457607. URL: <https://doi.org/10.1145/3457607>.
- [42] Tim Meinhardt et al. “TrackFormer: Multi-Object Tracking with Transformers”. In: *IEEE Conf. Comput. Vis. Pattern Recog.* 2022.
- [43] Volodymyr Mnih et al. “Human-level control through deep reinforcement learning”. In: *Nature* 518.7540 (Feb. 2015), pp. 529–533. ISSN: 1476-4687. DOI: 10.1038/nature14236. URL: <https://doi.org/10.1038/nature14236>.
- [44] Allen Newell and Herbert A. Simon. “Computer science as empirical inquiry: symbols and search”. In: *Communications of the ACM* 19.3 (Mar. 1976), pp. 113–126. DOI: 10.1145/1283920.1283930. URL: <http://doi.acm.org/10.1145/1283920.1283930>.
- [45] Alec Radford et al. “Language Models are Unsupervised Multitask Learners”. In: (2019).
- [46] Aditya Ramesh et al. *Hierarchical Text-Conditional Image Generation with CLIP Latents*. 2022. DOI: 10.48550/ARXIV.2204.06125. URL: <https://arxiv.org/abs/2204.06125>.
- [47] Aditya Ramesh et al. “Zero-Shot Text-to-Image Generation”. In: *Proceedings of the 38th International Conference on Machine Learning*. Ed. by Marina Meila and Tong Zhang. Vol. 139. Proceedings of Machine Learning Research. PMLR, 18–24 Jul 2021, pp. 8821–8831. URL: <https://proceedings.mlr.press/v139/ramesh21a.html>.
- [48] Suman Ravuri et al. “Skilful precipitation nowcasting using deep generative models of radar”. In: *Nature* 597.7878 (Sept. 2021), pp. 672–677. ISSN: 1476-4687. DOI: 10.1038/s41586-021-03854-z. URL: <https://doi.org/10.1038/s41586-021-03854-z>.
- [49] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. “U-Net: Convolutional Networks for Biomedical Image Segmentation”. In: *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*. Ed. by Nassir Navab et al. Cham: Springer International Publishing, 2015, pp. 234–241. ISBN: 978-3-319-24574-4.
- [50] Kiersten M. Ruff and Rohit V. Pappu. “AlphaFold and Implications for Intrinsically Disordered Proteins”. In: *Journal of Molecular Biology* 433.20 (2021). From Protein Sequence to Structure at Warp Speed: How AlphaFold Impacts Biology, p. 167208. ISSN: 0022-2836. DOI: <https://doi.org/10.1016/j.jmb.2021.167208>. URL: <https://www.sciencedirect.com/science/article/pii/S0022283621004411>.
- [51] David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. “Learning Representations by Back-propagating Errors”. In: *Nature* 323.6088 (1986), pp. 533–536. DOI: 10.1038/323533a0. URL: <http://www.nature.com/articles/323533a0>.
- [52] Olga Russakovsky et al. “ImageNet Large Scale Visual Recognition Challenge”. In: *International Journal of Computer Vision* 115.3 (Dec. 2015), pp. 211–252. ISSN: 1573-1405. DOI: 10.1007/s11263-015-0816-y. URL: <https://doi.org/10.1007/s11263-015-0816-y>.
- [53] Alvaro Sanchez-Gonzalez et al. “Learning to Simulate Complex Physics with Graph Networks”. In: *Proceedings of the 37th International Conference on Machine Learning*. Ed. by Hal Daumé III and Aarti Singh. Vol. 119. Proceedings of Machine Learning Research. PMLR, 13–18 Jul 2020, pp. 8459–8468. URL: <https://proceedings.mlr.press/v119/sanchez-gonzalez20a.html>.
- [54] John Schulman et al. *Proximal Policy Optimization Algorithms*. 2017. DOI: 10.48550/ARXIV.1707.06347. URL: <https://arxiv.org/abs/1707.06347>.
- [55] Andrew W. Senior et al. “Improved protein structure prediction using potentials from deep learning”. In: *Nature* 577.7792 (Jan. 2020), pp. 706–710. ISSN: 1476-4687. DOI: 10.1038/s41586-019-1923-7. URL: <https://doi.org/10.1038/s41586-019-1923-7>.

- [56] David Silver et al. “A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play”. In: *Science* 362.6419 (2018), pp. 1140–1144. DOI: 10.1126/science.aar6404. eprint: <https://www.science.org/doi/pdf/10.1126/science.aar6404>. URL: <https://www.science.org/doi/abs/10.1126/science.aar6404>.
- [57] David Silver et al. “Mastering the Game of Go with Deep Neural Networks and Tree Search”. In: *Nature* 529.7587 (jan 2016), pp. 484–489. ISSN: 0028-0836. DOI: 10.1038/nature16961.
- [58] David Silver et al. “Reward is enough”. In: *Artificial Intelligence* 299 (2021), p. 103535. ISSN: 0004-3702. DOI: <https://doi.org/10.1016/j.artint.2021.103535>. URL: <https://www.sciencedirect.com/science/article/pii/S0004370221000862>.
- [59] Emma Strubell, Ananya Ganesh, and Andrew McCallum. “Energy and Policy Considerations for Deep Learning in NLP”. In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*. Florence, Italy: Association for Computational Linguistics, July 2019, pp. 3645–3650. DOI: 10.18653/v1/P19-1355. URL: <https://aclanthology.org/P19-1355>.
- [60] Emma Strubell, Ananya Ganesh, and Andrew McCallum. “Energy and Policy Considerations for Modern Deep Learning Research”. In: *Proceedings of the AAAI Conference on Artificial Intelligence* 34.09 (Apr. 2020), pp. 13693–13696. DOI: 10.1609/aaai.v34i09.7123. URL: <https://ojs.aaai.org/index.php/AAAI/article/view/7123>.
- [61] Chi Sun et al. “How to Fine-Tune BERT for Text Classification?”. In: *Chinese Computational Linguistics*. Ed. by Maosong Sun et al. Cham: Springer International Publishing, 2019, pp. 194–206. ISBN: 978-3-030-32381-3.
- [62] Yi Sun et al. *NSP-BERT: A Prompt-based Zero-Shot Learner Through an Original Pre-training Task-Next Sentence Prediction*. 2021. arXiv: 2109.03564 [cs.CL].
- [63] A. M. Turing. “Computing Machinery and Intelligence”. In: *Mind* LIX.236 (Oct. 1950), pp. 433–460. ISSN: 0026-4423. DOI: 10.1093/mind/LIX.236.433. URL: <https://doi.org/10.1093/mind/LIX.236.433>.
- [64] Ashish Vaswani et al. *Attention Is All You Need*. 2017. DOI: 10.48550/ARXIV.1706.03762. URL: <https://arxiv.org/abs/1706.03762>.
- [65] Oriol Vinyals et al. “Grandmaster level in StarCraft II using multi-agent reinforcement learning”. In: *Nature* 575.7782 (2019), pp. 350–354.
- [66] Mark Weber et al. “STEP: Segmenting and Tracking Every Pixel”. In: *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 2)*. 2021. URL: <https://openreview.net/forum?id=0uQIr4XA77f>.
- [67] Joseph Weizenbaum. “ELIZA—a Computer Program for the Study of Natural Language Communication between Man and Machine”. In: *Commun. ACM* 9.1 (Jan. 1966), pp. 36–45. ISSN: 0001-0782. DOI: 10.1145/365153.365168. URL: <https://doi.org/10.1145/365153.365168>.
- [68] Xinhua Wu and Xiujie Liu. “Building crack identification and total quality management method based on deep learning”. In: *Pattern Recognition Letters* 145 (2021), pp. 225–231. ISSN: 0167-8655. DOI: <https://doi.org/10.1016/j.patrec.2021.01.034>. URL: <https://www.sciencedirect.com/science/article/pii/S0167865521000532>.
- [69] Haoran Xu, Benjamin Van Durme, and Kenton Murray. “BERT, mBERT, or BiBERT? A Study on Contextualized Embeddings for Neural Machine Translation”. In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Online and Punta Cana, Dominican Republic: Association for Computational Linguistics, Nov. 2021, pp. 6663–6675. DOI: 10.18653/v1/2021.emnlp-main.534. URL: <https://aclanthology.org/2021.emnlp-main.534>.
- [70] Zhuosheng Zhang, Junjie Yang, and Hai Zhao. *Retrospective Reader for Machine Reading Comprehension*. 2020. DOI: 10.48550/ARXIV.2001.09694. URL: <https://arxiv.org/abs/2001.09694>.